



Swing Trader · Control Centre

One settings file, edited from six panels and read left to right. Each panel saves its own part, **Run the test** on the Scoreboard spends every saved setting at once, and the Scoreboard then ranks the new fit against every fit on record. The Ledger keeps the money.

CONTROL CENTRE SHEET

05-research/scripts/control_centre.sh → 127.0.0.1:8787
offline copy 01-dashboards/control-centre.html
State of 24 September 2026

A1 Data

what is read, and what counts as a win

A1 Market and basket

Which prices the model studies. Values shown are as shipped; each tool has its own Save.

Choose Market

Market crypto Binance archives, or equity for adjusted Alpaca bars

Timeframe 4h slice bar size, 5 minutes to 1 day, for coins or stocks

Choose Basket

Quick pick all a named list, such as majors or large caps

Coins or stocks empty typed names, which win over the quick pick

History to load 40,000 candles read, newest first; 0 reads all

Quick picks are majors, scalp-eight, btc-eth, large-caps and sector-funds.

The 4h slice is a 25 MB test cut with no bitcoin in it. Judge a named coin on the full 4h panel.

A1 Filter and label

Choose Filter

Volume floor 30m USDT traded a day

Volatility band 1.5–7.1% usual daily move

Minimum history 157 days age of the asset

Choose Ranking

Rank by none order assets each bar

Keep third all top, bottom or all

Fold pass rate 60% folds that must pay

Choose Label

Outcome barrier win or loss, or three-way

Take-profit, stop +2 / -1 ATR in usual moves

Horizon 12 bars two days of 4h bars

Break-even band 0.20% three-way only

SCREENED BARS STILL HAVE PRICES The filter removes about 60% of a stock's days as entries. A trade's exit can still land on one of those days.

CHECK THE LABEL FIRST The +2/-1 barrier is hit on 0.313 of bars where 0.333 breaks even, so it loses before any model is fitted.

TUNING LOOP

one pass per change, repeated

1 Change one thing

On one panel, change one setting and press that tool's Save. Two changes at once cannot be told apart.

2 Run the test

On C1. The line under the button is the exact command it will run with every saved setting.

3 Read three bars

U2 below 1 on the blind period, overfit ratio at or below 1.1, and 60% of folds beating a constant guess.

4 Compare

The record line says whether this run beat the best on file. A gain under 0.0021 is noise.

5 Keep or revert

Keep the setting, or press Load best as defaults to return to the best fit on record.

Retuning rules

Never tune on the blind period, since a setting chosen by looking at it has spent it. Trust only walk-forward regimes, judge on a full blind year, and rerun on names the model has not seen.

A2 Indicators

the columns the model may see

A2 Feature families

Choose Features

Families none ticked nothing ticked offers every family

Also include empty exact column names, added

Leave out empty exact column names, removed last

At most 0 above 0 keeps the best by AUC alone

PREFIX WHAT THE COLUMNS MEASURE

f_wc_ f_hr_ price over day-length and short windows

f_ta_ f_ta_pta_ oscillators such as Stochastic, CCI, ADX

f_tl_ TA-Lib extras, candle patterns

f_st_ f_mst_ Supertrend, fixed and adaptive

f_btc_ strength against bitcoin

f_4h_ f_d1_ f_w1_ slower bars, joined without lookahead

f_flow_ buyers against sellers

f_rg_ f_ms_ regime state and microstructure

A2 Indicator engines

Four engines each reduce a chart to a stance at every bar. Nothing runs on this panel.

Choose MACD

Spans 12 / 26 / 9 fast, slow, signal

Noise band 0.5 a cross must clear it

Confirm bars 1 bars it must hold

Choose Averages

Fast, slow 20 / 50 bars in each average

Choose Fibonacci

Lookback 240 bars searched for a swing

Ignore swings 0 smaller than this share of price

MACD is the 12-bar less the 26-bar exponential average, against its own 9-bar average.

Choose Confluence

Score to fire 2 engines that must agree

Candle lasts 3 bars a candle signal counts

B1 Variables

which columns earn a place

B1 Variable screen

An elastic net on the training window pushes every coefficient toward zero; the columns still standing are the survivors.

Choose Screen

Screen first off run the screen before the fit

Fit on survivors on fit on survivors only; off reports the screen and keeps every column

L1 mix 1.0 1 is lasso, 0 is ridge

Penalty rule 1se loosest penalty within one standard error of the best

Rows 25,000 sample the screen is fitted on

Folds 10 for choosing the penalty only

Draw intervals on 95% intervals from an unpenalised refit

Screen variables Univariate screen

Screen variables runs the elastic net. Univariate screen scores each column alone against a model with no columns.

B1 What to read

90 columns offered penalty survivors with intervals

fitted on the training window only, never the blind period

FAMILY, 4H ALONE STEADY USED

vs bitcoin .019 .83 .0069

pandas-ta .049 1.00 .0023

adaptive .008 .60 .0004

Supertrend

TA-Lib extras .012 .75 .0003

8 September. alone is AUC on its own less 0.5, steady the share of five folds keeping its sign, used how much the score worsens when the column is shuffled. Strength against bitcoin leads by a factor of three.

Read which columns survive and whether their intervals cross zero. A column whose interval crosses zero has not shown which way it pushes.

MEASURE THE COST FIRST Run with SCREEN FIRST off, then on. The difference in the blind score on the Scoreboard is what the screen bought or cost.

MOST COLUMNS CARRY LITTLE On the 4-hour panel 16 of 90 barely moved the score when shuffled, and 14 of 90 changed sign between folds.

B2 Training

how the past is split

B2 Blind and folds

training window, cut into folds blind period

Blind days 365 held back, scored once

Embargo bars 0 0 uses the label horizon as the gap

Choose Blind Period

Choose Resampling

Regime expanding how the training window is cut

Folds 3 pieces of the training window

Purge bars 0 dropped at each fold edge

Repeats 10 repeated k-fold only

Training share 0.75 Monte Carlo only

Bootstrap samples 25 bootstrap only

Choose winner by best or oneSE, the simplest close to best

Compare regimes

Compare regimes scores one model under all seven regimes against the same blind period, so the gap between claim and result is the regime's own.

B2 What to read

Only expanding and rolling keep time in order. The other five put later bars in training and earlier ones in test.

CLAIMED IN TRAINING AGAINST THE BLIND YEAR, 24 SEP

random regimes

walk-forward

blind year

Per cent a trade after cost, the most each claimed. The regime changes the claim, never the fitted model.

BEST MINUS WORST MODEL, BY LENGTH OF BLIND STRETCH

30 days

60 days

90 days

180 days

365 days

Points a trade after cost, median of six setups, eight models. Short stretches spread the models apart by chance.

SHORT BLIND PERIODS PICK AT RANDOM On one fit, 30-day stretches of the blind year crowned 3 to 7 different winners out of 12. Judge on the full year and on several stretches.

C1 Scoreboard

fit, score, tune, compare

C1 Models and tuning

Choose Model

Models all six two logistic, forest, three boosters

Class weight balanced none measured better calibrated

Overfit ratio cap 1.1 reject above it

Choose Settings

Hyperparameters per model 9 forest, 12 LightGBM, 9 HistGBM

How hard to tune 0 3 quick, 10 thorough, 0 uses the grid

Choose Grid Search

Grid search over empty name one model to tune it

Grid key=v1,v2 the values to try

Choose Calibration

Calibrate on Platt, isotonic, or both

Held-out fraction 0.2 end of training, never blind

Run the test Compare models

Assess models Audit split Score edge

Trend-life search

C1 Three bars

U2 below 1 beat a constant guess, blind

ratio ≤ 1.1 not memorised

60% of folds beat a constant

FITS ON RECORD, BY THE BARS THEY PASSED

every fit

ratio at or under 1.1

U2 below 1

both

U2 is the model's error on the blind period divided by the error of always guessing the average. The ratio is held-out error over training error.

574 fits on record

117 beat a constant and passed

0.980 best U2, forest, 21 Sep

The record opens the panel and says whether the last run beat the best so far. Three-way lists fits scored in money, where the best is +1.226% a trade after cost, 3 coins, 16 September, with 12 distinct results from 42 fits.

NOISE FLOOR 0.0021 The bench's own randomness over three repeats. A change in U2 smaller than that is not an improvement.

C2 Ledger

the book and the history of runs

C2 The paper book

Nothing is configured here except the figures. Money gives each holding in dollars, paid, worth now and its effect on the whole account.

\$102,206 from \$100,000 on 18 Aug

+2.2% book, against +0.1% market

+\$4,481 open positions

-\$2,265 58 finished trades

6 / 6 weekly rebalances

book +2.2% +0.1%

Daily closes of the paper account against the same \$100,000 in the S&P 500 fund, SPY.

Marked 23 September 2026, 12:17 New York. Fifty shares at about 1.8% each, trading cost 13.2 basis points against 5 to 10 assumed.

PAPER MONEY Every order path stops unless the live switch reads exactly true. Never short, never borrow, 5% most in one name, 10% least in cash.

C2 Runs and figures

Live book is the run calendar and the 18 milestones where the design or the data changed. Assessment is every run by its own best fit, 125 so far, so a setting can be set against the one it replaced.

Choose Figures

Figures two candles, MACD, reliability, importance and more

Overlays none EMA 200, Supertrend, entries, exits

Symbol, bars first, 400 what is charted

Theme light or dark

MONEY COLUMN WHAT IT SAYS

paid, worth now dollars at the fill and at the last mark

gain or loss the difference, in dollars

share of book the holding's slice of the account

effect on account its move times its share; the column that adds up

EVERY ROW NAMES ITS RECORD

Each fit and each run links to the file it was read from under 04-outputs/AA-evals, so a number on the page can be checked against its source.