



Swing Trader · After-Fee Workflow

For a reader who trades nothing and has trained no model. Five columns, one per stage, read left to right. **Is there a rule that clears the fee, on data it has never seen?** Every stage answers a piece of that, and no number here was written without opening its file.

WORKFLOW CHEAT SHEET

seamusrobertmurphy/swing-trader
Crypto NO-GO · Equities live on paper
State of 7 September 2026

A What counts as winning the bar every idea is held to

A1 The fee wall

THE WHOLE STORY

Every trade pays a toll in and another out, so a faster clock pays more often. The 5-minute frame has an edge before fees; the toll is six times it.



the argument, one line kept = move - round trip - slippage
allow if kept $\approx 1.5\%$


0.20	0.15	1.5
% assumed round trip	% best case round trip	% the least worth having

THIS IS THE WALL A basis point is a hundredth of one per cent. The best crypto edge is five to twelve of them *smaller* than the fifteen to twenty a round trip costs. Nothing tried has closed it.

A2 What counts as a win

THREE BARRIERS

A win is a decision someone makes, not a fact waiting to be found. Three lines are drawn from the moment of buying: a profit to aim at, a loss to bail out at, and a deadline. Whichever the price touches first is the answer.



Both lines are set in units of the coin's own daily travel, so a calm coin and a violent one are judged on equal terms.

A TARGET CAN LOSE BEFORE THE MODEL SPEAKS

On the 4-hour panel the profit line is reached first on **254** bars in every thousand. Aiming twice as far up as down, you need **333** to come out level. The arithmetic is losing before anything is fitted, so the first thing to check is always the target's own record, not the model's.

A3 Passing the bar


FOLDS, NOT TOTALS

A grand total can be carried by one lucky year. So history is cut into half-year blocks and the edge has to show up in at least **60 per cent** of them, on two questions at once.

QUESTION **WHAT IT ASKS**

Did it pay did the picks make money after fees

Did it beat did they beat the average name in the same block



NEVER READ THE AVERAGE ALONE Those five blocks average out to exactly what a steady 24.9 would, and a steady 24.9 would be a usable model. These are not.

TRAPS THAT COST DAYS

every one of these happened here, and each was found late

A number quoted from memory

Three figures on the previous version of this sheet were wrong because they were recalled rather than read. Open the file that produced the number, every time, including your own.

Rendering is not checking

A page can build, exit cleanly, and report 0 of 6 cycles where there are 3. Compare the new render against the last one before believing it.

The same error, twice

The overfit ratio was written upside down, corrected on 7 September in one file, and found the same day still inverted in a second. A model reported as stable at 0.902 was really 1.109, which the rule rejects.

B Where the data comes from sources, screening, ranking

B1 Where the numbers come from

TWO ARCHIVES

Two filing cabinets, not two live feeds. Both were downloaded once and are read off disk, so a result rebuilt next year comes out the same.

Crypto. Binance's own monthly files of open, high, low, close and volume, one zip per coin per month, checksum-verified.

Shares. Alpaca daily bars, split and dividend adjusted. Of 12,571 live US companies, **2,660** clear \$20m a day above \$3 a share.

PANEL	ROWS	NAMES	INPUTS
5 min	5.53M	11	96
1 hour	1.64M	47	61
4 hour	3.94M	567	90
1 day	0.52M	540	104
shares 1d	5.00M	2,547	96

DEAD COINS ARE KEPT **612** pairs ever traded against **487** alive today. Build the list from today's exchange and you delete every coin that failed, which makes any result an upper bound. The 1-hour panel is still the old 47-coin build.

B2 Getting into the universe

ALL FOUR OR OUT

Four doors, and a coin has to pass all four before it is even considered. The test is re-run on every bar using only what was known on that day, because a filter applied with hindsight is not a filter.

DOOR	BAR IT HAS TO CLEAR
Trading enough	$\geq 30m$ USD in 24h
Moving enough	ATR 2.5–12% of price
Cheap to cross	$\leq 0.5\%$ estimated
Old enough	≥ 120 days lived

ATR, the average true range, is how far price usually travels in a day, as a share of price. The band shown is daily; a 4-hour bar wants 1.0 to 4.9%.

The **spread**, the gap between what buyers offer and sellers ask, is paid both ways. The archive holds no order book, so it is estimated from each day's high and low and gated at 0.5%.

RANK WIDE, BUY NARROW The whole market is scored, then only coins that actually carried the edge on unseen data are allowed through at the moment of buying.

B3 Rank, do not predict

THE SURVIVING IDEA

Two different questions, and only one of them has ever produced anything. **Will this coin go up?** failed on every frame. **Which of these coins is strongest right now?** did not.

EVERY BAR, SORT THE WHOLE MARKET

TOP THIRD	MIDDLE	BOTTOM THIRD
hold these	ignore	never sold short

Of 42 ways of sorting tried, 24 June 2026:

25 pointed the same way on new data
33 had the top third beat the market

And still losing, because the market itself lost 0.382% a trade.

TWO CLAIMS, KEPT APART The sorting works and has never been disproved. The one way of *trading it* that was built, waiting for bitcoin to trend up first, held in 27 per cent of blocks against the 60 the rule wants, and was killed. Only the second sentence is settled. Untried: holding longer to pay the toll less often, and running the same sort on shares, where the market drifts up.

C What the model reads signals, and which ones earn their place

C1 What the model reads

FOUR GROUPS

THE COIN'S OWN PRICE HISTORY		
day/length windows	short windows	chart indicators
WHICH WAY THE TREND POINTS		
Supertrend	adaptive version	calm or wild
WHAT THE REST OF THE MARKET IS DOING		
bitcoin, leading	slower bars above	crowding
WHO WAS IN A HURRY TO TRADE		
buyers against sellers	order-book texture	

Every input reads only bars that have already closed, and every one is a ratio, so it means the same on a \$90,000 coin as on a coin worth a tenth of a cent.

TWENTY-FOUR NUMBERS DO THE WORK OF NINETY Asked how many bars a trend has left, the 90 built inputs miss by **24.213** bars on the 4-hour panel. Twenty-four plain past returns, nothing built at all, miss by **24.247**. Guessing the average misses by 26.477. The two real answers sit a thirty-fourth of a bar apart, so they hold the same information and a memory-based model has no gap left to close.

C2 Which inputs earn their place

3 OF 90 FAMILIES

Ninety inputs were built and most are not read. Each scored three ways, on the training window only.

alone how far its own ranking sits from a coin flip
steady in how many of five stretches it pointed the same way
used how much worse the model scores with that column scrambled

FAMILY	N	ALONE	STEADY	USED
vs bitcoin	7	.019	.83	.0069
pandas-ta	7	.049	1.00	.0023
weekly view	5	.017	.72	.0021
Supertrend	7	.031	.89	.0006
adaptive Supertrend	6	.008	.60	.0004
TA-Lib extras	15	.012	.75	.0003

THE MARKET, NOT THE COIN The strongest family by three times is a coin's strength against bitcoin, led by **F.BTC.MOM.168**, its week-long move against the market. It is the only family not built from the coin's own price, and the one the model leans on hardest.

DEAD WEIGHT Eight of the ten columns never read are candlestick patterns. **16 OF 90** inputs barely move the score when scrambled; **14 OF 90** flip direction between stretches of market.

D How it is fitted splits, tuning, training

D1 Splitting time

NEVER SHUFFLED

Shuffle the rows and you put next Tuesday in the study pile, then test on last Monday. Prices carry over from one bar to the next, so a shuffle lets the future leak backwards and every score comes out flatter.



The yellow sliver is the **gap**. Without it the last rows studied still await outcomes that land inside the test window, the same leak in a different hat.

ONCE MEANS ONCE The final year is opened at the end and scored a single time. A setting chosen by looking at it has spent it, and the year cannot be made unseen again.


AS CONFIGURED

PIECE	SETTING
Held back	the final 365 days
Studied	all history before it
Gap at the cut	8 days (48 4h bars)

D2 What tuning showed

SMALLER, NOT BIGGER

Nine settings, each tried on three separate stretches of the 4-hour panel, all asked the same question: how many bars until the trend turns?



The miss on studied bars nearly halves while the miss on unseen bars gets worse. All nine settings fail the 1.1 rule, and the one that came top was the least powerful in the grid.

WHEN BEATS HOW Two experiments agree. The nine settings above span **1.366** bars end to end; a separate test of three ways of feeding the model spanned **0.094**. Move one model from a calm stretch of market to a rough one and it shifts **3.83**. Which months you test on moves the score further than any choice you make about the model.

E What came out assessment, results, the live book

E1 How a model is judged

TWO DEFECTS

A model answers with a probability, so the score is how far it sat from what happened.

ratio unseen miss divided by studied miss

MODEL	STUDIED	UNSEEN	RATIO
HistGBM	.2619	.4818	1.840
LightGBM	.2806	.4826	1.720
Random forest	.4526	.4840	1.069
Logistic	.4855	.4909	1.011


1-hour panel, 6 Sep 2026. A ratio near 1 generalises. Above 1.1 it has memorised, and is thrown out.

TWO FAULTS, BOTH IN THE JUDGING The winner was chosen on the unseen miss alone, never the ratio, crowning HistGBM at **1.840** where the rule rejects past 1.1. And every model treats rare wins as common, pushing its probabilities away from what this score measures: on 4-hour bars it misses by **0.4844** where answering "25.4 per cent, always" misses **0.4353**. Nothing checks whether a stated 30 per cent happens 30 per cent of the time.

E2 Where it stands

THE SCOREBOARD

Kept per crypto trade after fees, on bars never studied. A year of work moved the line and never crossed zero.



MOST LOSE BEFORE THE FEE Fee stripped out: daily **-1.150%** a trade, 4-hour **-0.182%**. Only 5-minute is ahead, **+0.033%**, which the toll wipes out six times over.


THE ONE THAT WORKS

Twelve-month share momentum. Buy what rose most over the past year, skip the latest month, hold. Over 116 months it beat the average share by **1.006** points a month, **2.41** times its own wobble, where two is the bar. It paid in 85 per cent of blocks and beat the market in 80.

Closed: 1-hour direction, a coin flip; sharper entries, 29% of blocks; the breadth gate, 27%.

E3 The live book

PAPER, REAL PRICES



Fourteen daily closes against holdings at Friday's prices, \$96,916.58 from \$100,000. The plan expects an 8.2% swing; a 3.1% dip is four tenths of it.

49 shares held at 1.8% each	3/6 clean weekly cycles done	7.6 bp a fill mid-session
---------------------------------------	--	-------------------------------------

weekly, never at the opening bell
alpaca_trade.py rebalance --execute

THIS IS FAKE MONEY A paper account at real prices; every order path stops unless it reads the exact word true. Fixed in code: never short, never borrow, never options, never add to a loser, 5% most in one name, 10% least in cash.

A paper account is not a mark

On 7 September, a market holiday, the broker reported \$98,362 while the holdings at Friday's real closes came to \$96,917. Value the positions yourself before reporting a number.