

Evaluation scores

date - the day the evaluation was run.

evaluation type - which kind of evaluation. Head-to-head: several models trained on the same train/test split and compared. Later types: walk-forward backtest, tuning sweep, stability check.

dataset - rows and feature count used (e.g. 26,762r / 32f).

best model - the model kept, chosen by the highest buy-class precision on the test set.

test AUC - area under the ROC curve on the out-of-sample test set: the chance the model scores a real buy above a non-buy. 0.50 = no skill (a coin flip), 1.00 = perfect ranking; it does not depend on a threshold.

precision - of the days the model called a buy (probability ≥ 0.5), the share that were genuine buys: $TP / (TP + FP)$. "When it says buy, how often is it right."

base rate - the share of all test days that were genuine buys, $P(\text{label}=1)$. The precision you would get by blindly calling everything a buy: the baseline to beat.

lift (x) - precision divided by base rate. How many times better than the blind baseline. 1.0x = no better than guessing; above 1 = adding value. Verify it yourself: precision / base rate.

verdict - GO only if precision clearly beats the base rate and AUC clears 0.55; otherwise NO-GO.

record - links to the full per-run report (markdown for the numbers, HTML for the charts).

A day is labelled a "buy" (label = 1) if its close gains +10% before falling -5% within the next 20 days; otherwise 0. Each per-run record also reports Keller Metric 1 (precision and recall at the 60/40 trading threshold) and Metric 3 (AUC split by volatility regime).

date	evaluation type	dataset	best model	test AUC	precision	base rate	lift (x)	verdict
2026-06-20	head-to-head	26,762r / 32f	LightGBM	0.520	0.325	0.305	1.07x	NO-GO